



< Previous	✓	✓	✓	✓									Next >
------------	---	---	---	---	--	--	--	--	--	--	--	--	--------

## Part 4: How To

[Bookmark this page](#)

### Part 4: How To

How to actually DO threading and multiprocessing:

- `threading` module
- `multiprocessing` module

Parallel programming can be hard!

If your problem can be solved sequentially, consider the costs and benefits before going parallel.

### Parallelization Strategy for Performance

1. Break problem down into chunks.
2. Execute chunks in parallel.
3. Reassemble output of chunks into a result.

- Not every problem is parallelizable.
- There is an optimal number of threads for each problem in each environment, so make it tunable.
- Working concurrently opens up synchronization issues.
- Methods for synchronizing threads:
  - locks.
  - queues.
  - signaling/messaging mechanisms.

### The mechanics: how do you use threads and/or processes

Python provides the *threading* and *multiprocessing* modules to facilitate concurrency.

They have similar APIs – so you can use them in similar ways.

Key points:

- There is no Python thread scheduler, it is up to the host OS. Yes, these are “true” threads.
- Works well for I/O bound problems, can use literally thousands of threads

- works well for I/O bound problems, can use literally thousands of threads.
- Limit CPU-bound processing to C extensions (that release the GIL).
- Do not use for CPU bound problems, will go slower than no threads, especially on multiple cores!!! (see David Beazley's talk referenced above).

Starting threads is relatively simple, but there are many potential issues.

We already talked about shared data, this can lead to a "race condition".

- May produce slightly different results every run.
- May just flake out mysteriously every once in a while.
- May run fine when testing, but fail when run on: - a slower system. - a heavily loaded system. - a larger dataset.
- Thus you *must* synchronize threads!

### Example: A CPU bound problem

Numerically integrate the function

$$y = x^2$$

from 0 to 10.

<http://www.wolframalpha.com/input/?i=x%5E2>

### Solution

Parallel execution example

Consider the following code in: [integrate.py](#)

```
def f(x):  
    return x**2  
  
def integrate(f, a, b, N):
```

```
s = 0
dx = (b-a)/N
for i in range(N):
    s += f(a+i*dx)
return s * dx
```

We can do better than this!



© 2024 University of Washington | Seattle, WA. All rights reserved.

[Help Center](#) [Contact Us](#) [Privacy](#) [Terms](#)

Built on  by  

edX, Open edX and the edX and Open edX logos are trademarks or registered trademarks of edX Inc.